

English Accent Detection: A Comparison of Audio, Spectrogram and Text Classification Method using Transformers

Colab Link: <https://colab.research.google.com/drive/1V59vfCSLv6DxCXBrqy1M9k081DuEAcjB?usp=sharing>

Candidate Numbers: 31948, 24692, 24155

Date: 29 April 2024

Audio Array Classification

In [35]:

```
%%capture
! pip install datasets evaluate
! pip install accelerate
! pip install wandb
```

In [2]:

```
import os, sys
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

In [37]:

```
import pandas as pd
import numpy as np
import os
```

In []:

```
from datasets import Dataset, Audio, load_dataset
from sklearn.model_selection import train_test_split
import pyarrow.parquet as pq
from transformers import Wav2Vec2Processor
import random
from sklearn.utils.class_weight import compute_class_weight
```

In [5]:

```
from transformers import AutoModelForAudioClassification, TrainingArguments, Trainer, AutoFeatureExtractor
import evaluate
from datasets import Dataset, Audio
import torch
```

In [6]:

```
from transformers import EarlyStoppingCallback, IntervalStrategy
```

In [7]:

```
from datasets import load_from_disk
```

In [8]:

```
import wandb
wandb.login()
wandb.init(project="accents_classification")
```

```
wandb: Logging into wandb.ai. (Learn how to deploy a W&B server locally: https://wandb.me/wandb-server)
wandb: You can find your API key in your browser here: https://wandb.ai/authorize
wandb: Paste an API key from your profile and hit enter, or press ctrl+c to quit:
```

.....

```
wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc
wandb: Currently logged in as: yiwe (yiweizhou). Use `wandb login --relogin` to force relogin
```

Tracking run with wandb version 0.16.6

Run data is saved locally in `/content/wandb/run-20240428_211525-5bhqlbh6`

Syncing run [revived-bee-5](#) to [Weights & Biases \(docs\)](#)

View project at https://wandb.ai/yiweizhou/accents_classification

View run at https://wandb.ai/yiweizhou/accents_classification/runs/5bhqlbh6

Out[8]:

Display W&B run

In []:

```
# Change directory to data location
os.chdir("/content/drive/My Drive/ST311")
```

In []:

```
# Load data
df0 = pq.read_table(source= '0000.parquet').to_pandas()
df1 = pq.read_table(source= '0001.parquet').to_pandas()
df2 = pq.read_table(source= '0002.parquet').to_pandas()
df3 = pq.read_table(source= '0003.parquet').to_pandas()
df4 = pq.read_table(source= '0004.parquet').to_pandas()
df5 = pq.read_table(source= '0005.parquet').to_pandas()
df6 = pq.read_table(source= '0006.parquet').to_pandas()
df7 = pq.read_table(source= '0007.parquet').to_pandas()
data = pd.concat([df0,df1,df2,df3,df4,df5,df6,df7])
data.reset_index(drop=True, inplace=True)
```

In []:

```
data = data.drop(['path'], axis=1)
data = data.drop(['up_votes'], axis=1)
data = data.drop(['down_votes'], axis=1)
data = data.drop(['age'], axis=1)
data = data.drop(['gender'], axis=1)

#train test split
train_data, test_data = train_test_split(data, test_size=0.2, random_state=42)
```

In []:

```
data.value_counts()
```

In [14]:

```
print(len(data), len(train_data), len(test_data))
```

7678 6142 1536

In []:

```
test_data.iloc[1529]
```

Out[]:

accents

United States English

```
accents          United States English
sentence          They are on the ground.
audio            {'array': [2.2737367544323206e-12, 9.094947017...
Name: 7482, dtype: object
```

```
In [ ]:
```

```
test_data['accents'].value_counts()
```

```
Out[ ]:
```

```
accents
United States English          670
India and South Asia (India, Pakistan, Sri Lanka)    225
England English                203
Filipino                      136
Australian English             135
Canadian English              100
Irish English                  39
Scottish English              28
Name: count, dtype: int64
```

```
In [ ]:
```

```
train_data['accents'].value_counts()
```

```
Out[ ]:
```

```
accents
United States English          2577
India and South Asia (India, Pakistan, Sri Lanka)    963
England English                891
Australian English             528
Filipino                      471
Canadian English              395
Irish English                  226
Scottish English              91
Name: count, dtype: int64
```

```
In [ ]:
```

```
test_data.iloc[0,2]
```

```
Out[ ]:
```

```
{'array': array([ 0.00000000e+00,  2.18278728e-11,  1.45519152e-11, ...,
                 -4.17522733e-06, -4.00739464e-06,  1.34157199e-06]),
 'path': './accents/datasets/common_voice_16_1/en/clips/common_voice_en_20275899.mp3',
 'sampling_rate': 16000}
```

```
In [ ]:
```

```
weights = compute_class_weight(class_weight = "balanced", classes = np.unique(data['accen
ts']), y = train_data['accents'])
print(weights)
```

```
[1.45407197  1.94367089  0.86167228  1.63004246  0.79724818  3.39712389
 8.43681319  0.29792394]
```

```
In [ ]:
```

```
train_ds = Dataset.from_pandas(train_data)
test_ds = Dataset.from_pandas(test_data)
```

```
In [ ]:
```

```
train_ds = train_ds.class_encode_column("accents")
test_ds = test_ds.class_encode_column("accents")
```

20% 80% split\ accents column类型转换了

Audio Classification without Spectrogram

In []:

```
train_ds = train_ds.remove_columns("sentence")
test_ds = test_ds.remove_columns("sentence")
```

In []:

```
# labels = accents_ds["train"].features["accents"].names
# label2id, id2label = dict(), dict()
# for i, label in enumerate(labels):
#     label2id[label] = str(i)
#     id2label[str(i)] = label
```

In []:

```
label2id = {'Australian English': '0',
            'Canadian English': '1',
            'England English': '2',
            'Filipino': '3',
            'India and South Asia (India, Pakistan, Sri Lanka)': '4',
            'Irish English': '5',
            'Scottish English': '6',
            'United States English': '7'}
id2label = {'0': 'Australian English',
            '1': 'Canadian English',
            '2': 'England English',
            '3': 'Filipino',
            '4': 'India and South Asia (India, Pakistan, Sri Lanka)',
            '5': 'Irish English',
            '6': 'Scottish English',
            '7': 'United States English'}
```

In [10]:

```
feature_extractor = AutoFeatureExtractor.from_pretrained("facebook/wav2vec2-base")
```

```
/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_token.py:89: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it as secret in your Google Colab and restart your session.
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
  warnings.warn(
```

```
/usr/local/lib/python3.10/dist-packages/transformers/configuration_utils.py:363: UserWarning: Passing `gradient_checkpointing` to a config initialization is deprecated and will be removed in v5 Transformers. Using `model.gradient_checkpointing_enable()` instead, or if you are using the `Trainer` API, pass `gradient_checkpointing=True` in your `TrainingArguments`.
  warnings.warn(
```

In [11]:

```
feature_extractor
```

Out[11]:

```
Wav2Vec2FeatureExtractor {
  "do_normalize": true,
  "feature_extractor_type": "Wav2Vec2FeatureExtractor",
  "feature_size": 1,
  "padding_side": "right",
  "padding_value": 0.0,
  "return_attention_mask": false,
  "sampling_rate": 16000
}
```

```
In [ ]:
```

```
def transform_audio(example):  
    array = np.array(example['audio']['array'], dtype=np.float32)  
    sampling_rate = example['audio']['sampling_rate'] # Already set correctly, just refo  
rencing it  
    return {'audio': {'array': array, 'sampling_rate': sampling_rate}}
```

```
In [ ]:
```

```
train_ds = train_ds.map(transform_audio)  
test_ds = test_ds.map(transform_audio)
```

```
In [ ]:
```

```
train_ds = train_ds.cast_column("audio", Audio(sampling_rate=16_000))  
test_ds = test_ds.cast_column("audio", Audio(sampling_rate=16_000))
```

```
In [ ]:
```

```
def preprocess_function(examples):  
    audio_arrays = [x["array"] for x in examples["audio"]]  
    inputs = feature_extractor(  
        audio_arrays, sampling_rate=feature_extractor.sampling_rate, max_length=160000, t  
runcation=True  
    )  
  
    return inputs  
  
# Apply preprocessing  
train_ds = train_ds.map(preprocess_function, remove_columns=["audio"], batched=True)  
test_ds = test_ds.map(preprocess_function, remove_columns=["audio"], batched=True)  
  
# Rename 'accents' column to 'label'  
train_ds = train_ds.rename_column("accents", "label")  
test_ds = test_ds.rename_column("accents", "label")
```

```
In [ ]:
```

```
train_ds.save_to_disk("train_23")  
test_ds.save_to_disk("test_23")
```

```
In [ ]:
```

```
train_ds = load_from_disk("train_23")  
test_ds = load_from_disk("test_23")
```

```
In [ ]:
```

```
# accuracy = evaluate.load("accuracy")  
# f1 = evaluate.load("f1")  
# def compute_metrics(eval_pred):  
#     predictions = np.argmax(eval_pred.predictions, axis=1)  
#     return {  
#         'accuracy': accuracy.compute(predictions=predictions, references=eval_pred.label  
_ids),  
#         'f1': f1.compute(predictions=predictions, references=eval_pred.label_ids, avera  
ge='weighted')  
#     }
```

```
In [ ]:
```

```
accuracy = evaluate.load("accuracy")  
f1 = evaluate.load("f1")  
precision = evaluate.load("precision")  
recall = evaluate.load("recall")  
mcc = evaluate.load("matthews_correlation")
```

```
def compute_metrics(eval_pred):
    predictions = np.argmax(eval_pred.predictions, axis=1)
    references = eval_pred.label_ids
    return {
        'accuracy': accuracy.compute(predictions=predictions, references=references)['accuracy'],
        'f1_weighted': f1.compute(predictions=predictions, references=references, average='weighted')['f1'],
        'precision_weighted': precision.compute(predictions=predictions, references=references, average='weighted')['precision'],
        'recall_weighted': recall.compute(predictions=predictions, references=references, average='weighted')['recall'],
        'mcc': mcc.compute(predictions=predictions, references=references)['matthews_correlation']
    }
```

In []:

```
num_labels = 8
model = AutoModelForAudioClassification.from_pretrained(
    "facebook/wav2vec2-base", num_labels=num_labels, label2id=label2id, id2label=id2label
)
```

/usr/local/lib/python3.10/dist-packages/transformers/configuration_utils.py:363: UserWarning: Passing `gradient_checkpointing` to a config initialization is deprecated and will be removed in v5 Transformers. Using `model.gradient_checkpointing_enable()` instead, or if you are using the `Trainer` API, pass `gradient_checkpointing=True` in your `TrainingArguments`.

```
warnings.warn(
Some weights of Wav2Vec2ForSequenceClassification were not initialized from the model checkpoint at facebook/wav2vec2-base and are newly initialized: ['classifier.bias', 'classifier.weight', 'projector.bias', 'projector.weight', 'wav2vec2.encoder.pos_conv_embed.conv.parametrizations.weight.original0', 'wav2vec2.encoder.pos_conv_embed.conv.parametrizations.weight.original1']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
```

In []:

```
ordered_weights = [1.4475867269984917,
1.9388888888888889,
0.8772851919561243,
1.581136738056013,
0.8078703703703703,
3.621698113207547,
8.065126050420169,
0.2955805358792732]
```

In []:

```
classes = list(id2label.values())
class_weights = dict(zip(classes, weights))
print(class_weights)
ordered_weights = [class_weights[x] for x in id2label.values()]
```

In []:

```
wandb.init(project="accents_classification_28", name = "run_2")
```

Finishing last run (ID:8cklz8mv) before initializing another...

Run history:

eval/accuracy




eval/f1_weighted



eval/loss





Run summary:

eval/accuracy	0.71354
eval/f1_weighted	0.7255
eval/loss	1.65135
eval/mcc	0.67784
eval/precision_weighted	0.82566
eval/recall_weighted	0.71354
eval/runtime	71.2861
eval/samples_per_second	21.547
eval/steps_per_second	0.673
train/epoch	1.45833
train/global_step	70
train/grad_norm	2.82574
train/learning_rate	2e-05
train/loss	1.3193

View run **run_1** at: https://wandb.ai/yiweizhou/accents_classification_28/runs/8cklz8mv

View project at: https://wandb.ai/yiweizhou/accents_classification_28

Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: `/content/wandb/run-20240428_144952-8cklz8mv/logs`

Successfully finished last run (ID:8cklz8mv). Initializing new run:

Tracking run with wandb version 0.16.6

Run data is saved locally in `/content/wandb/run-20240428_151625-4w4gl03o`

Syncing run **run_2** to [Weights & Biases \(docs\)](#)

View project at https://wandb.ai/yiweizhou/accents_classification_28

View run at https://wandb.ai/yiweizhou/accents_classification_28/runs/4w4gl03o

Out[]:

Display W&B run

In []:

```

training_args = TrainingArguments (
    output_dir="wav2vec-23",
    evaluation_strategy=IntervalStrategy.STEPS,
    eval_steps = 50,
    save_strategy=IntervalStrategy.STEPS,
    save_steps = 100,
    save_total_limit = 4,
    learning_rate=2e-5,
    per_device_train_batch_size=32,
    gradient_accumulation_steps=4,
    per_device_eval_batch_size=32,
    num_train_epochs=10,
    warmup_ratio=0.1,
    logging_steps=10,
    load_best_model_at_end=True,
    metric_for_best_model="accuracy",
    push_to_hub=False,
    report_to="wandb"
)

```

In []:

```

class WeightedTrainer(Trainer):
    def compute_loss(self, model, inputs, return_outputs=False):
        labels = inputs.pop("labels")
        # forward pass
        outputs = model(**inputs)
        logits = outputs.get("logits")
        # compute custom loss (suppose one has labels with different weights)
        loss_fct = torch.nn.CrossEntropyLoss(weight=torch.tensor(ordered_weights, device=
model.device).float())
        loss = loss_fct(logits.view(-1, self.model.config.num_labels), labels.view(-1))
        return (loss, outputs) if return_outputs else loss

```

In []:

```

trainer = WeightedTrainer(
    model=model,
    args=training_args,
    train_dataset=train_ds,
    eval_dataset=test_ds,
    tokenizer=feature_extractor,
    compute_metrics=compute_metrics,
    callbacks = [EarlyStoppingCallback(early_stopping_patience=2)]
)

```

In []:

```
trainer.evaluate()
```

██ [48/48 01:10]

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: Undefined MetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))

Out[]:

```

{'eval_loss': 2.081277847290039,
 'eval_accuracy': 0.095703125,
 'eval_f1_weighted': 0.11839429320686339,
 'eval_precision_weighted': 0.27227572133292505,
 'eval_recall_weighted': 0.095703125,
 'eval_mcc': -0.0022599274251548247,
 'eval_runtime': 73.9545,
 'eval_samples_per_second': 20.77,
 'eval_steps_per_second': 0.649}

```

In []:

```
trained = trainer.train()
```




[462/480 1:16:22 < 02:59, 0.10 it/s, Epoch 9.60/10]

Step	Training Loss	Validation Loss	Accuracy	F1 Weighted	Precision Weighted	Recall Weighted	Mcc
50	1.837200	1.764204	0.757812	0.746339	0.794171	0.757812	0.676719
100	1.001900	0.891448	0.973307	0.973553	0.975202	0.973307	0.965212
150	0.555400	0.492019	0.995443	0.995428	0.995512	0.995443	0.993934
200	0.342900	0.296086	0.994792	0.994799	0.994855	0.994792	0.993071
250	0.236300	0.204503	0.998047	0.998049	0.998065	0.998047	0.997397
300	0.181300	0.164038	0.997396	0.997394	0.997412	0.997396	0.996528
350	0.147500	0.132456	0.998698	0.998701	0.998714	0.998698	0.998265
400	0.132700	0.117824	0.998698	0.998701	0.998714	0.998698	0.998265
450	0.123200	0.112002	0.998047	0.998050	0.998062	0.998047	0.997397



[48/48 09:30]



[480/480 1:19:03, Epoch 10/10]

Step	Training Loss	Validation Loss	Accuracy	F1 Weighted	Precision Weighted	Recall Weighted	Mcc
50	1.837200	1.764204	0.757812	0.746339	0.794171	0.757812	0.676719
100	1.001900	0.891448	0.973307	0.973553	0.975202	0.973307	0.965212
150	0.555400	0.492019	0.995443	0.995428	0.995512	0.995443	0.993934
200	0.342900	0.296086	0.994792	0.994799	0.994855	0.994792	0.993071
250	0.236300	0.204503	0.998047	0.998049	0.998065	0.998047	0.997397
300	0.181300	0.164038	0.997396	0.997394	0.997412	0.997396	0.996528
350	0.147500	0.132456	0.998698	0.998701	0.998714	0.998698	0.998265
400	0.132700	0.117824	0.998698	0.998701	0.998714	0.998698	0.998265
450	0.123200	0.112002	0.998047	0.998050	0.998062	0.998047	0.997397

In []:

```
trainer.evaluate()
```



[48/48 01:09]

Out[]:

```
{'eval_loss': 0.11782387644052505,
 'eval_accuracy': 0.9986979166666666,
 'eval_f1_weighted': 0.9987012334284002,
 'eval_precision_weighted': 0.9987141927083334,
 'eval_recall_weighted': 0.9986979166666666,
 'eval_mcc': 0.9982646808828041,
 'eval_runtime': 72.4762,
 'eval_samples_per_second': 21.193,
 'eval_steps_per_second': 0.662,
 'epoch': 10.0}
```

In []:

```
trained
```

Out[]:

```
TrainOutput(global_step=240, training_loss=0.7800431281328202, metrics={'train_runtime': 3
625.1638, 'train_samples_per_second': 8.471, 'train_steps_per_second': 0.066, 'total_flos'
: 2.400507469299843e+18, 'train_loss': 0.7800431281328202, 'epoch': 5.0})
```

T= 5.1

```
In [ ]:
```

```
trainer.save_model("audio_27apr")
```

```
In [ ]:
```

```
trainer.evaluate()
```

```
In [4]:
```

```
#confusion matrix
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [ ]:
```

```
predictions = trainer.predict(test_ds)
y_pred = np.argmax(predictions.predictions, axis=1)
y_true = test_ds['label']
```

```
In [ ]:
```

```
y_true_names = [id2label[str(id)] for id in y_true]
y_pred_names = [id2label[str(id)] for id in y_pred]
```

```
In [ ]:
```

```
cm2 = confusion_matrix(y_true_names, y_pred_names, labels = list(id2label.values()))
```

```
In [ ]:
```

```
cm2
```

```
Out[ ]:
```

```
array([[135,  0,  0,  0,  0,  0,  0,  0],
       [  0, 100,  0,  0,  0,  0,  0,  0],
       [  0,  0, 202,  0,  0,  0,  0,  1],
       [  0,  0,  0, 136,  0,  0,  0,  0],
       [  0,  0,  0,  0, 225,  0,  0,  0],
       [  0,  0,  0,  0,  0, 39,  0,  0],
       [  0,  0,  0,  0,  0,  0, 28,  0],
       [  0,  0,  0,  0,  0,  1,  0, 669]])
```

```
In [6]:
```


```
cm = np.array([[135,  0,  0,  0,  0,  0,  0,  0],
               [  0, 100,  0,  0,  0,  0,  0,  0],
               [  0,  0, 202,  0,  0,  0,  0,  1],
               [  0,  0,  0, 136,  0,  0,  0,  0],
               [  0,  0,  0,  0, 225,  0,  0,  0],
               [  0,  0,  0,  0,  0, 39,  0,  0],
               [  0,  0,  0,  0,  0,  0, 28,  0],
               [  0,  0,  0,  0,  0,  1,  0, 669]])
```

```
In [34]:
```

```
plt.figure(figsize=(6, 2))
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('True (X) and False Labels (Y)', fontsize = 10)
ax = sns.heatmap(cm, annot=True, fmt='g', cmap='Blues')
```

```
# Saving the figure with a high DPI
plt.savefig('cm_audio.png', dpi=1000)
```





In []:

```
np.flatnonzero(y_true != y_pred)
```

Out[]:

```
array([ 469, 1529])
```

In []:

```
id2label.values()
```

Out[]:

```
dict_values(['Australian English', 'Canadian English', 'England English', 'Filipino', 'India and South Asia (India, Pakistan, Sri Lanka)', 'Irish English', 'Scottish English', 'United States English'])
```

In []:

```
id2label.values()
```

Out[]:

```
dict_values(['Australian English', 'Canadian English', 'England English', 'Filipino', 'India and South Asia (India, Pakistan, Sri Lanka)', 'Irish English', 'Scottish English', 'United States English'])
```

In []:

```
y_true[469]
```

Out[]:

```
2
```

In []:

```
y_pred[469]
```

Out[]:

```
7
```

In []:

```
y_true[1529]
```

Out[]:

```
7
```

In []:

```
y_pred[1529]
```

Out[]:

```
5
```

In []:

```
# pred = irish
test_data.iloc[[1529]]
```

Out[]:

	accents	sentence	audio
7482	United States English	They are on the ground.	{'array': [2.2737367544323206e-12, 9.094947017...

In []:

```
# pred = US
test_data.iloc[[469]]
```

Out[]:

	accents	sentence	audio
965	England English	As you wish.	{'array': [0.0, 0.0, 0.0, 0.0, 0.0, ...

In []:

```
# from numba import cuda
# device = cuda.get_current_device()
# device.reset()
```

In []:

```
# model_local = AutoModelForAudioClassification.from_pretrained('audio_27apr')
```

Spectrogram Classification

In []:

```
# !pip install librosa
# !pip install --upgrade transformers
# !pip install accelerate -U
```

In []:

```
import os
import torch
import torchvision
import torch.nn as nn
import matplotlib.pyplot as plt
from torchvision import datasets, transforms
from torch.utils.data import DataLoader, SubsetRandomSampler
import torch.optim as optim
import time
import numpy as np
import pandas as pd
from tqdm.autonotebook import tqdm
from PIL import Image
import IPython.display as ipd
import numpy as np
import librosa

cuda = torch.cuda.is_available()
```

In []:

```
import pyarrow.parquet as pq
```

```
df0 = pq.read_table(source= '0000.parquet').to_pandas()
df1 = pq.read_table(source= '0001.parquet').to_pandas()
df2 = pq.read_table(source= '0002.parquet').to_pandas()
df3 = pq.read_table(source= '0003.parquet').to_pandas()
df4 = pq.read_table(source= '0004.parquet').to_pandas()
df5 = pq.read_table(source= '0005.parquet').to_pandas()
df6 = pq.read_table(source= '0006.parquet').to_pandas()
df7 = pq.read_table(source= '0007.parquet').to_pandas()
```

Method 2: Add noise, change pitch and change speed

In []:

```
data = pd.concat([df0,df1,df2,df3,df4,df5,df6,df7])
data.reset_index(drop=True, inplace=True)
data = data.drop(['path'], axis=1)
data = data.drop(['up_votes'], axis=1)
data = data.drop(['down_votes'], axis=1)
data = data.drop(['age'], axis=1)
data = data.drop(['gender'], axis=1)

#train test split
from sklearn.model_selection import train_test_split
data, test_data = train_test_split(data, test_size=0.2, random_state=42)

data['accents'].value_counts()
```

Out[]:

```
United States English                2577
India and South Asia (India, Pakistan, Sri Lanka)    963
England English                      891
Australian English                   528
Filipino                           471
Canadian English                   395
Irish English                      226
Scottish English                    91
Name: accents, dtype: int64
```

In []:

```
test_data['accents'].value_counts()
```

Out[]:

```
United States English                670
India and South Asia (India, Pakistan, Sri Lanka)    225
England English                      203
Filipino                           136
Australian English                   135
Canadian English                   100
Irish English                      39
Scottish English                    28
Name: accents, dtype: int64
```

In []:

```
new_audio = []
for r in data['audio']:
    new_audio.append(r['array'])
data['audio'] = new_audio
```

In []:

```
us = data[data['accents'] == 'United States English']
us = us[:800]

indian = data[data['accents'] == 'India and South Asia (India, Pakistan, Sri Lanka)']
indian = indian[:800]

england = data[data['accents'] == 'England English']
```

```
england = england[:800]
```

```
filipino = data[data['accents'] == 'Filipino']  
filipino.reset_index(drop=True, inplace=True)
```

```
australian = data[data['accents'] == 'Australian English']  
australian.reset_index(drop=True, inplace=True)
```

```
scottish = data[data['accents'] == 'Scottish English']  
scottish.reset_index(drop=True, inplace=True)
```

```
irish = data[data['accents'] == 'Irish English']  
irish.reset_index(drop=True, inplace=True)
```

```
canadian = data[data['accents'] == 'Canadian English']  
canadian.reset_index(drop=True, inplace=True)
```

In []:

```
scottish_arrays = scottish['audio']  
irish_arrays = irish['audio']  
canadian_arrays = canadian['audio']  
australian_arrays = australian['audio']  
filipino_arrays = filipino['audio']
```

In []:

```
sampling_rate = 16000  # This is the sampling rate from your data  
  
# Function to add noise  
def add_noise(audio, noise_level=0.005):  
    noise = np.random.randn(len(audio))  
    augmented_audio = audio + noise_level * noise  
    return augmented_audio  
  
# Function to change pitch  
def change_pitch(audio, sampling_rate, n_steps=5):  
    return librosa.effects.pitch_shift(audio, sr=sampling_rate, n_steps=n_steps)  
  
# Function to change speed  
def change_speed(audio, speed_factor=1.5):  
    return librosa.effects.time_stretch(audio, rate=speed_factor)
```

In []:

```
scottish_augmented = []  
irish_augmented = []  
canadian_augmented = []  
australian_augmented = []  
filipino_augmented = []  
  
for i in range(len(scottish_arrays)):  
    audio = scottish_arrays[i]  
    noisy_audio = add_noise(audio)  
    pitched_audio = change_pitch(audio, sampling_rate)  
    speeded_audio = change_speed(audio)  
    scottish_augmented.append(noisy_audio)  
    scottish_augmented.append(pitched_audio)  
    scottish_augmented.append(speeded_audio)  
  
for i in range(190):  
    audio = irish_arrays[i]  
    noisy_audio = add_noise(audio)  
    pitched_audio = change_pitch(audio, sampling_rate)  
    speeded_audio = change_speed(audio)  
    irish_augmented.append(noisy_audio)  
    irish_augmented.append(pitched_audio)  
    irish_augmented.append(speeded_audio)  
  
for i in range(130):  
    audio = canadian_arrays[i]
```

```

noisy_audio = add_noise(audio)
pitched_audio = change_pitch(audio, sampling_rate)
speeded_audio = change_speed(audio)
canadian_augmented.append(noisy_audio)
canadian_augmented.append(pitched_audio)
canadian_augmented.append(speeded_audio)

for i in range(75):
    audio = australian_arrays[i]
    noisy_audio = add_noise(audio)
    pitched_audio = change_pitch(audio, sampling_rate)
    speeded_audio = change_speed(audio)
    australian_augmented.append(noisy_audio)
    australian_augmented.append(pitched_audio)
    australian_augmented.append(speeded_audio)

for i in range(80):
    audio = filipino_arrays[i]
    noisy_audio = add_noise(audio)
    pitched_audio = change_pitch(audio, sampling_rate)
    speeded_audio = change_speed(audio)
    filipino_augmented.append(noisy_audio)
    filipino_augmented.append(pitched_audio)
    filipino_augmented.append(speeded_audio)

```

In []:

```

scottish_augmented_sentence = []
irish_augmented_sentence = []
canadian_augmented_sentence = []
australian_augmented_sentence = []
filipino_augmented_sentence = []

for i in range(len(scottish)):
    scottish_augmented_sentence.append(scottish['sentence'][i])
    scottish_augmented_sentence.append(scottish['sentence'][i])
    scottish_augmented_sentence.append(scottish['sentence'][i])

for i in range(190):
    irish_augmented_sentence.append(irish['sentence'][i])
    irish_augmented_sentence.append(irish['sentence'][i])
    irish_augmented_sentence.append(irish['sentence'][i])

for i in range(130):
    canadian_augmented_sentence.append(canadian['sentence'][i])
    canadian_augmented_sentence.append(canadian['sentence'][i])
    canadian_augmented_sentence.append(canadian['sentence'][i])

for i in range(75):
    australian_augmented_sentence.append(australian['sentence'][i])
    australian_augmented_sentence.append(australian['sentence'][i])
    australian_augmented_sentence.append(australian['sentence'][i])

for i in range(80):
    filipino_augmented_sentence.append(filipino['sentence'][i])
    filipino_augmented_sentence.append(filipino['sentence'][i])
    filipino_augmented_sentence.append(filipino['sentence'][i])

```

In []:

```

scottish_augmented_df = pd.DataFrame({'accents': 'Scottish English', 'sentence': scottish_
augmented_sentence, 'audio': scottish_augmented})
irish_augmented_df = pd.DataFrame({'accents': 'Irish English', 'sentence': irish_augmented
_sentence, 'audio': irish_augmented})
canadian_augmented_df = pd.DataFrame({'accents': 'Canadian English', 'sentence': canadian_
augmented_sentence, 'audio': canadian_augmented})
australian_augmented_df = pd.DataFrame({'accents': 'Australian English', 'sentence': austr
alian_augmented_sentence, 'audio': australian_augmented})
filipino_augmented_df = pd.DataFrame({'accents': 'Filipino', 'sentence': filipino_augment
ed_sentence, 'audio': filipino_augmented})

scottish = pd.concat([scottish, scottish_augmented_df])

```

```

scottish.reset_index(drop=True, inplace=True)
irish = pd.concat([irish, irish_augmented_df])
irish.reset_index(drop=True, inplace=True)
canadian = pd.concat([canadian, canadian_augmented_df])
canadian.reset_index(drop=True, inplace=True)
australian = pd.concat([australian, australian_augmented_df])
australian.reset_index(drop=True, inplace=True)
filipino = pd.concat([filipino, filipino_augmented_df])
filipino.reset_index(drop=True, inplace=True)

```

In []:

```

data_augmented = pd.concat([us, indian, england, scottish, irish, canadian, australian, filipino])

```

In []:

```

#second round augmentation for scottish
scottish_augmented = []
for i in range(120, 240):
    audio = scottish['audio'][i]
    noisy_audio = add_noise(audio)
    pitched_audio = change_pitch(audio, sampling_rate)
    speeded_audio = change_speed(audio)
    scottish_augmented.append(noisy_audio)
    scottish_augmented.append(pitched_audio)
    scottish_augmented.append(speeded_audio)

scottish_augmented_sentence = []
for i in range(120, 240):
    scottish_augmented_sentence.append(scottish['sentence'][i])
    scottish_augmented_sentence.append(scottish['sentence'][i])
    scottish_augmented_sentence.append(scottish['sentence'][i])

scottish_augmented_df = pd.DataFrame({'accents': 'Scottish English', 'sentence': scottish_augmented_sentence, 'audio': scottish_augmented})
data_augmented = pd.concat([data_augmented, scottish_augmented_df])

```

In []:

```

data_augmented['accents'].value_counts()

```

Out[]:

United States English	800
India and South Asia (India, Pakistan, Sri Lanka)	800
England English	800
Irish English	796
Canadian English	785
Australian English	753
Scottish English	724
Filipino	711

Name: accents, dtype: int64

In []:

```

test_audio = []
for audio in test_data['audio']:
    test_audio.append(audio['array'])
test_data['audio'] = test_audio

```

In []:

```

test_data['accents'].value_counts()
test_data.reset_index(drop=True, inplace=True)

```

In []:

```

data_augmented.reset_index(drop=True, inplace=True)

```

Audio to Image spectra

Audio to image spectro

- ViT (Vision Transformer)

In []:

```
import torch
import torchvision.models as models
import torch.nn as nn
import matplotlib.pyplot as plt
from torchvision import datasets, transforms
from torch.utils.data import DataLoader, SubsetRandomSampler
import torch.optim as optim # SGD
import time
import numpy as np
from tqdm.autonotebook import tqdm #adds a progress bar to track training
import pandas as pd
```

In []:

```
#turn audio array to spectrogram
def audio_to_melspectrogram(audio, sr = sampling_rate):
    melspec = librosa.feature.melspectrogram(y=audio, sr = sr, n_mels=128)
    melspec_db = librosa.power_to_db(melspec, ref=np.max)
    return melspec_db
```

In []:

```
#turn all audio arrays to melspectrograms
melspectrograms = []
for audio in data_augmented['audio']:
    melspec = audio_to_melspectrogram(audio)
    melspectrograms.append(melspec)
```

In []:

```
melspectrograms_test = []
for audio in test_data['audio']:
    melspec = audio_to_melspectrogram(audio)
    melspectrograms_test.append(melspec)
```

In []:

```
#pip install opencv-python
```

Collecting opencv-python

Downloading opencv_python-4.9.0.80-cp37-abi3-macosx_10_16_x86_64.whl.metadata (20 kB)
Requirement already satisfied: numpy>=1.17.0 in /Users/dongzhehu/opt/anaconda3/lib/python3.8/site-packages (from opencv-python) (1.24.4)
Downloading opencv_python-4.9.0.80-cp37-abi3-macosx_10_16_x86_64.whl (55.7 MB)
55.7/55.7 MB 15.8 MB/s eta 0:00:000:0100:01

Installing collected packages: opencv-python

Successfully installed opencv-python-4.9.0.80

Note: you may need to restart the kernel to use updated packages.

In []:

```
import cv2
import numpy as np
def resize_images(images, size=(224, 224)):
    resized_images = []
    for img in images:
        resized_image = cv2.resize(img, size, interpolation=cv2.INTER_CUBIC) # Using INTER_CUBIC for better resizing quality
        resized_images.append(resized_image)
    return np.array(resized_images)

# # Call the resize function
resized_images = resize_images(melspectrograms)
resized_images_test = resize_images(melspectrograms_test)
```

In []:

```
#normalize the images
normalized_images = abs(resized_images / 255)
normalized_images_test = abs(resized_images_test / 255)
```

In []:

```
#3 channel images
final_images = []
for img in normalized_images:
    final_image = np.stack((img,)*3, axis=-1)
    final_images.append(final_image)
final_images = np.array(final_images)

final_images_test = []
for img in normalized_images_test:
    final_image = np.stack((img,)*3, axis=-1)
    final_images_test.append(final_image)
final_images_test = np.array(final_images_test)
```

In []:

```
#output images to a folder
import os
import cv2
import numpy as np
from PIL import Image
import matplotlib.pyplot as plt

def save_images(images, path):
    for i in range(len(images)):
        image = images[i]
        image = Image.fromarray((image * 255).astype(np.uint8))
        image.save(f'{path}/image_{i}.png')

save_images(final_images, 'train_images')
save_images(final_images_test, 'test_images')
```

```
-----
NameError                                Traceback (most recent call last)
/Users/dongzhehu/Documents/ST311 Final/parquet.ipynb Cell 30 in 1
      1 <a href='vscode-notebook-cell:/Users/dongzhehu/Documents/ST311%20Final/parquet.ipynb#
Y540sZmlsZQ%3D%3D?line=10'>11</a>         image = Image.fromarray((image * 255).astype(np.
uint8))
      2 <a href='vscode-notebook-cell:/Users/dongzhehu/Documents/ST311%20Final/parquet.ipynb#
Y540sZmlsZQ%3D%3D?line=11'>12</a>         image.save(f'{path}/image_{i}.png')
---> <a href='vscode-notebook-cell:/Users/dongzhehu/Documents/ST311%20Final/parquet.ipynb#
Y540sZmlsZQ%3D%3D?line=13'>14</a> save_images(final_images, 'train_images')
      4 <a href='vscode-notebook-cell:/Users/dongzhehu/Documents/ST311%20Final/parquet.ipynb#
Y540sZmlsZQ%3D%3D?line=14'>15</a> save_images(final_images_test, 'test_images')
```

NameError: name 'final_images' is not defined

In []:

```
data_augmented['accents'] = data_augmented['accents'].astype('category')
test_data['accents'] = test_data['accents'].astype('category')
```

In []:

```
lables = data_augmented['accents']
lables_test = test_data['accents']
#output the labels to a csv file
lables.to_csv('text_labels.csv', index=False)
lables_test.to_csv('text_labels_test.csv', index=False)
```

In []:

```
# #melspectrogram back to audio
# def melspectrogram_to_audio(spec_img):
```

```
# mel_inverted = librosa.feature.inverse.mel_to_audio(spec_img, sr=sampling_rate, n_iter=32)
# return mel_inverted

# audio_test = melspectrogram_to_audio(melspec_df['melspectrogram'][0])
# ipd.Audio(audio_test, rate=sampling_rate)
```

Load Local Data

In []:

```
# #unzip file
# import zipfile
# with zipfile.ZipFile('train_images.zip', 'r') as zip_ref:
#     zip_ref.extractall('train_images')
# with zipfile.ZipFile('test_images.zip', 'r') as zip_ref:
#     zip_ref.extractall('test_images')
```

In []:

```
from datasets import Dataset, Image
#load all images
train_images_path = []
for i in range(0, 6169):
    train_images_path.append(f'train_images/image_{i}.png')

test_images_path = []
for i in range(0, 1536):
    test_images_path.append(f'test_images/image_{i}.png')
```

In []:

```
train_dataset = Dataset.from_dict({"image": train_images_path}).cast_column("image", Image())
test_dataset = Dataset.from_dict({"image": test_images_path}).cast_column("image", Image())
```

In []:

```
#remove column name
import pandas as pd
train_labels = pd.read_csv('text_labels.csv')
test_labels = pd.read_csv('text_labels_test.csv')
```

In []:

```
#replace each label with number
train_labels['accents'] = train_labels['accents'].replace({'United States English': 0, 'India and South Asia (India, Pakistan, Sri Lanka)': 1, 'England English': 2, 'Scottish English': 3, 'Irish English': 4, 'Canadian English': 5, 'Australian English': 6, 'Filipino': 7})
test_labels['accents'] = test_labels['accents'].replace({'United States English': 0, 'India and South Asia (India, Pakistan, Sri Lanka)': 1, 'England English': 2, 'Scottish English': 3, 'Irish English': 4, 'Canadian English': 5, 'Australian English': 6, 'Filipino': 7})
```

In []:

```
#add the labels to the dataset
train_dataset = train_dataset.add_column("labels", train_labels['accents'])
test_dataset = test_dataset.add_column("labels", test_labels['accents'])
```

In []:

In []:

```
from transformers import ViTImageProcessor
```

```
model_name_or_path = 'google/vit-base-patch16-224-in21k'
processor = ViTImageProcessor.from_pretrained(model_name_or_path)
```

In []:

```
#processing
def process_example(example):
    inputs = processor(example['image'], return_tensors='pt')
    inputs['labels'] = example['labels']
    return inputs

process_example(train_dataset[0])
process_example(test_dataset[0])
```

Out[]:

```
{'pixel_values': tensor([[[[-0.6000, -0.6549, -0.6863, ..., -0.6078, -0.5843, -0.5451],
  [-0.5922, -0.6392, -0.6706, ..., -0.5922, -0.5765, -0.5451],
  [-0.5686, -0.6157, -0.6392, ..., -0.5765, -0.5686, -0.5451],
  ...,
  [-0.3725, -0.3804, -0.3725, ..., -0.3725, -0.3725, -0.3725],
  [-0.3725, -0.3804, -0.3725, ..., -0.3725, -0.3725, -0.3725],
  [-0.3725, -0.3725, -0.3804, ..., -0.3725, -0.3725, -0.3725]]],

  [[[-0.6000, -0.6549, -0.6863, ..., -0.6078, -0.5843, -0.5451],
  [-0.5922, -0.6392, -0.6706, ..., -0.5922, -0.5765, -0.5451],
  [-0.5686, -0.6157, -0.6392, ..., -0.5765, -0.5686, -0.5451],
  ...,
  [-0.3725, -0.3804, -0.3725, ..., -0.3725, -0.3725, -0.3725],
  [-0.3725, -0.3804, -0.3725, ..., -0.3725, -0.3725, -0.3725],
  [-0.3725, -0.3725, -0.3804, ..., -0.3725, -0.3725, -0.3725]]],

  [[[-0.6000, -0.6549, -0.6863, ..., -0.6078, -0.5843, -0.5451],
  [-0.5922, -0.6392, -0.6706, ..., -0.5922, -0.5765, -0.5451],
  [-0.5686, -0.6157, -0.6392, ..., -0.5765, -0.5686, -0.5451],
  ...,
  [-0.3725, -0.3804, -0.3725, ..., -0.3725, -0.3725, -0.3725],
  [-0.3725, -0.3804, -0.3725, ..., -0.3725, -0.3725, -0.3725],
  [-0.3725, -0.3725, -0.3804, ..., -0.3725, -0.3725, -0.3725]]]]), 'labels': 5}
```

In []:

```
def transform(example_batch):
    # Take a list of PIL images and turn them to pixel values
    inputs = processor([x for x in example_batch['image']], return_tensors='pt')

    # Don't forget to include the labels!
    inputs['labels'] = example_batch['labels']
    return inputs
```

In []:

```
prepared_ds = train_dataset.with_transform(transform)
prepared_ds_test = test_dataset.with_transform(transform)
```

In []:

```
#data collator
import torch

def collate_fn(batch):
    return {
        'pixel_values': torch.stack([x['pixel_values'] for x in batch]),
        'labels': torch.tensor([x['labels'] for x in batch])
    }

#evaluation metrics
import numpy as np
from datasets import load_metric
```

```
metric = load_metric("accuracy")
def compute_metrics(p):
    return metric.compute(predictions=np.argmax(p.predictions, axis=1), references=p.labels_ids)
```

Model

In []:

```
import wandb
wandb.login()
wandb.init(project='accent-classification')
```

wandb: Logging into wandb.ai. (Learn how to deploy a W&B server locally: <https://wandb.me/wandb-server>)
wandb: You can find your API key in your browser here: <https://wandb.ai/authorize>
wandb: Paste an API key from your profile and hit enter, or press ctrl+c to quit:**wandb**: Appending key for api.wandb.ai to your netrc file: /root/.netrc
wandb: Currently logged in as: **alexhudongzhe** (**aleksss**). Use `wandb login --relogin` to force relogin

wandb version 0.16.6 is available! To upgrade, please run: \$ pip install wandb --upgrade

Tracking run with wandb version 0.13.4

Run data is saved locally in /notebooks/wandb/run-20240429_150221-2puargng

Syncing run [azure-yogurt-3](#) to [Weights & Biases](#) ([docs](#))

Out[]:

Display W&B run

In []:

```
from transformers import ViTForImageClassification

labels = list(set(train_dataset['labels']))

model = ViTForImageClassification.from_pretrained(
    model_name_or_path,
    num_labels=len(labels),
    id2label={str(i): c for i, c in enumerate(labels)},
    label2id={c: str(i) for i, c in enumerate(labels)}
)
```

Some weights of ViTForImageClassification were not initialized from the model checkpoint at google/vit-base-patch16-224-in21k and are newly initialized: ['classifier.bias', 'classifier.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In []:


```
#arguments
from transformers import TrainingArguments
```

```
training_args = TrainingArguments(
    output_dir="./vit-base-beans",
    per_device_train_batch_size=32,
    evaluation_strategy="epoch",
    save_strategy="epoch",
    num_train_epochs=3,
    fp16=True,
    learning_rate=2e-4,
    save_total_limit=2,
    remove_unused_columns=False,
    push_to_hub=False,
```

```
predictions = trainer.predict(prepared_ds_test)
y_pred = np.argmax(predictions.predictions, axis=1)
y_true = test_labels
cm = confusion_matrix(y_true, y_pred)
```

Out[]:

<AxesSubplot: >



In []:

```
cm = ([[662, 2, 2, 0, 0, 0, 2, 2],
       [ 2, 223, 0, 0, 0, 0, 0, 0],
       [ 2, 0, 200, 0, 0, 1, 0, 0],
       [ 0, 0, 0, 28, 0, 0, 0, 0],
       [ 0, 0, 0, 0, 39, 0, 0, 0],
       [ 0, 0, 0, 0, 0, 100, 0, 0],
       [ 0, 0, 0, 0, 0, 0, 135, 0],
       [ 2, 0, 0, 0, 0, 0, 0, 134]])
```

```
cm2 = ([[459, 2, 47, 2, 152]])
```

```
cm3 = [
    [25, 2, 21, 32, 6, 1, 0, 48],
    [21, 12, 14, 17, 11, 0, 0, 25],
    [31, 12, 30, 42, 15, 3, 0, 70],
    [23, 5, 23, 26, 13, 2, 0, 44],
    [27, 3, 47, 47, 32, 5, 0, 64],
```

```

[6, 1, 3, 9, 2, 3, 0, 15],
[0, 0, 4, 8, 2, 0, 0, 8],
[137, 30, 103, 122, 55, 4, 1, 218]
]

cm3_reversed = [
[218, 1, 4, 55, 122, 103, 30, 137],
[8, 0, 0, 2, 8, 4, 0, 0],
[15, 0, 3, 2, 9, 3, 1, 6],
[64, 0, 5, 32, 47, 47, 3, 27],
[44, 0, 2, 13, 26, 23, 5, 23],
[70, 0, 3, 15, 42, 30, 12, 31],
[25, 0, 0, 11, 17, 14, 12, 21],
[48, 0, 1, 6, 32, 21, 2, 25]
]

cm4 = ([[669, 0, 0, 0, 1, 0, 0, 0],
[ 0, 225, 0, 0, 0, 0, 0, 0],
[ 1, 0, 202, 0, 0, 0, 0, 0],
[ 0, 0, 0, 28, 0, 0, 0, 0],
[ 0, 0, 0, 0, 39, 0, 0, 0],
[ 0, 0, 0, 0, 0, 100, 0, 0],
[ 0, 0, 0, 0, 0, 0, 135, 0],
[ 0, 0, 0, 0, 0, 0, 0, 136]])

```

In []:

```


plt.figure(figsize=(7, 3))
sns.heatmap(cm3_reversed, annot=True, fmt='g', cmap='Blues')
#x-axis labels

plt.xlabel('Predicted labels')
plt.ylabel('True labels')

```

Out[]:

Text(58.22222222222214, 0.5, 'True labels')



Extra Test

In []:

```

# load us audio
extra = pq.read_table(source= 'extra.parquet').to_pandas()

```

In []:

```

extra_audio = []
for r in extra['audio']:
    extra_audio.append(r['array'])

```



```

extra_melspectrograms = []
for audio in extra_audio:
    melspec = audio_to_melspectrogram(audio)
    extra_melspectrograms.append(melspec)

extra_resized_images = resize_images(extra_melspectrograms)
extra_normalized_images = abs(extra_resized_images / 255)
extra_final_images = []
for img in extra_normalized_images:
    final_image = np.stack((img,)*3, axis=-1)
    extra_final_images.append(final_image)
extra_final_images = np.array(extra_final_images)

```

In []:

```

#output images to a folder
def save_images(images, path):
    for i in range(len(images)):
        image = images[i]
        image = Image.fromarray((image * 255).astype(np.uint8))
        image.save(f'{path}/image_{i}.png')

save_images(extra_final_images, 'extra_images')

```

In []:

```

#save label
extra_labels = extra['accents']
#output the labels to a csv file
extra_labels.to_csv('extra_labels.csv', index=False)
#replace with numner
extra_labels = extra_labels.replace({'United States English': 0, 'India and South Asia (I
ndia, Pakistan, Sri Lanka)': 1, 'England English': 2, 'Scottish English': 3, 'Irish Englis
h': 4, 'Canadian English': 5, 'Australian English': 6, 'Filipino': 7})

```

In []:

```

#load images randomly
from datasets import Dataset, Image
extra_images_path = []
for i in range(0, 698):
    extra_images_path.append(f'extra_images/image_{i}.png')

extra_dataset = Dataset.from_dict({"image": extra_images_path}).cast_column("image", Imag
e())

```

In []:

```

#first 1000 labels
extra_labels = extra_labels[:1000]

```

In []:

```

extra_dataset = extra_dataset.add_column("labels", extra_labels)

```

In []:

```

#predict
prepared_extra_ds = extra_dataset.with_transform(transform)
predictions = trainer.predict(prepared_extra_ds)

```

In []:

```

predictions = np.argmax(predictions.predictions, axis=1)

```

In []:

```

predictions

```

In []:

```
#accuracy
from sklearn.metrics import accuracy_score
accuracy_score(extra_labels, predictions)
```

Out[]:

0.7091690544412608

In []:

```
#count unique in predictions
unique, counts = np.unique(predictions, return_counts=True)
dict(zip(unique, counts))
```

Out[]:

{0: 495, 2: 2, 3: 47, 5: 2, 7: 152}


In []:


```
#confusion matrix
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

cm = confusion_matrix(extra_labels, predictions)
plt.figure(figsize=(10, 10))
sns.heatmap(cm, annot=True, fmt='g', cmap='Blues')
#change x labels
plt.xticks(np.arange(5), ['United States English', 'England English', 'Scottish English',
'Canadian English', 'Filipino'], rotation=45)
```

Out[]:

```
(<matplotlib.axis.XTick at 0x7f92fe0cce50>,
<matplotlib.axis.XTick at 0x7f92fe0cce20>,
<matplotlib.axis.XTick at 0x7f92fe0ad880>,
<matplotlib.axis.XTick at 0x7f92fe0ba280>,
<matplotlib.axis.XTick at 0x7f92fe0bad30>],
[Text(0, 0, 'United States English'),
Text(1, 0, 'England English'),
Text(2, 0, 'Scottish English'),
Text(3, 0, 'Canadian English'),
Text(4, 0, 'Filipino')])
```





Extra 2

In []:

```
extra2 = pq.read_table(source= 'extra2.parquet').to_pandas()
```

In []:

```
extra_audio = []
for r in extra2['audio']:
    extra_audio.append(r['array'])

extra_melspectrograms = []
for audio in extra_audio:
    melspec = audio_to_melspectrogram(audio)
    extra_melspectrograms.append(melspec)

extra_resized_images = resize_images(extra_melspectrograms)
extra_normalized_images = abs(extra_resized_images / 255)
extra_final_images = []
for img in extra_normalized_images:
    final_image = np.stack((img,)*3, axis=-1)
    extra_final_images.append(final_image)
extra_final_images = np.array(extra_final_images)
```

In []:

```
#output images to a folder
from PIL import Image
def save_images(images, path):
    for i in range(len(images)):
        image = images[i]
        image = Image.fromarray((image * 255).astype(np.uint8))
        image.save(f'{path}/image_{i}.png')

save_images(extra_final_images, 'extra2_images')
```

In []:

```
extra2_labels = extra2['accents']
#output the labels to a csv file
extra2_labels.to_csv('extra2_labels.csv', index=False)
#replace with numner
extra2_labels = extra2_labels.replace({'United States English': 0, 'India and South Asia
(India, Pakistan, Sri Lanka)': 1, 'England English': 2, 'Scottish English': 3, 'Irish Engl
```

```
ish': 4, 'Canadian English': 5, 'Australian English': 6, 'Filipino': 7})
```

In []:

```
#load images randomly
from datasets import Dataset, Image
extra2_images_path = []
for i in range(0, 697):
    extra2_images_path.append(f'extra2_images/image_{i}.png')

extra2_dataset = Dataset.from_dict({"image": extra2_images_path}).cast_column("image", Image())
extra2_dataset = extra2_dataset.add_column("labels", extra2_labels)
```

In []:

```
#predict
prepared_extra2_ds = extra2_dataset.with_transform(transform)
predictions = trainer.predict(prepared_extra2_ds)
```

In []:

```
#accuracy
predictions = np.argmax(predictions.predictions, axis=1)
accuracy_score(extra2_labels, predictions)
```

Out[]:

0.975609756097561

Text Classification (Speech Transcript)

In []:

```
import os
import torch
import torchvision
import torch.nn as nn
import matplotlib.pyplot as plt
from torchvision import datasets, transforms
from torch.utils.data import DataLoader, SubsetRandomSampler, Subset
import torch.optim as optim
import time
import numpy as np
import pandas as pd
from tqdm.autonotebook import tqdm
from PIL import Image
import librosa
from datasets import Dataset, Audio
from sklearn.model_selection import train_test_split
from sklearn.utils.class_weight import compute_class_weight
```

```
cuda = torch.cuda.is_available()
```

Import the data

Data Source:

https://huggingface.co/datasets/WillHeld/accented_common_voice/tree/refs%2Fconvert%2Fparquet/default/parti
[train](#)

Text Transformer 1, PERT: <https://huggingface.co/google-bert/bert-base-uncased>

Due to the size and format of the original data file parquet, we decided to download the data and import it from our own laptop. The file link can be found above.

In []:

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

In []:

```
import pyarrow.parquet as pq

df0 = pq.read_table(source= '/content/drive/MyDrive/year3/ST311/0000.parquet').to_pandas()
df1 = pq.read_table(source= '/content/drive/MyDrive/year3/ST311/0001.parquet').to_pandas()
df2 = pq.read_table(source= '/content/drive/MyDrive/year3/ST311/0002.parquet').to_pandas()
df3 = pq.read_table(source= '/content/drive/MyDrive/year3/ST311/0003.parquet').to_pandas()
df4 = pq.read_table(source= '/content/drive/MyDrive/year3/ST311/0004.parquet').to_pandas()
df5 = pq.read_table(source= '/content/drive/MyDrive/year3/ST311/0005.parquet').to_pandas()
df6 = pq.read_table(source= '/content/drive/MyDrive/year3/ST311/0006.parquet').to_pandas()
df7 = pq.read_table(source= '/content/drive/MyDrive/year3/ST311/0007.parquet').to_pandas()
```

In []:

```
data = pd.concat([df0,df1,df2,df3,df4,df5,df6,df7])
data.reset_index(drop=True, inplace=True)
```

In []:

data

Out[]:

	path	accents	sentence	up_votes	down_votes	age	gender	audio
0	common_voice_en_23608720.mp3	India and South Asia (India, Pakistan, Sri Lanka)	The liner notes were written by Mark Paytress ...	2.0	1.0	twenties	female	{'array': [-1.7462298274040222e-10, -4.2200554...
1	common_voice_en_23608721.mp3	India and South Asia (India, Pakistan, Sri Lanka)	Microsoft has not made public statements about...	2.0	0.0	twenties	female	{'array': [5.093170329928398e-11, 7.3669070843...
2	common_voice_en_23608722.mp3	India and South Asia (India, Pakistan, Sri Lanka)	It indicates a letter that must be typed and m...	2.0	0.0	twenties	female	{'array': [0.0, -2.9103830456733704e-11, 4.365...
3	common_voice_en_23608754.mp3	India and South Asia (India, Pakistan, Sri Lanka)	Nasseri was born in the Anglo-Persian Oil Comp...	2.0	0.0	twenties	female	{'array': [5.4569682106375694e-11, -7.27595761...
		India						

	path	accents and South	He was sentence also an	up_votes	down_votes	age	gender	audio
4	common_voice_en_23608756.mp3	Asia (India, Pakistan, Sri Lanka)	ordained minister and led vario...	2.0	1.0	twenties	female	{'array': [-1.74622983e-10, 1.05218151e-05, 1.69793075e-05, ...]}
...
7673	common_voice_en_25122812.mp3	United States English	A Game Gear game by the same name was publishe...	2.0	0.0			{'array': [-2.4556356947869062e-11, -7.2759576...]}
7674	common_voice_en_25122824.mp3	United States English	He is the only President of Latvia to die in o...	2.0	0.0			{'array': [-1.0913936421275139e-11, -2.1827872...]}
7675	common_voice_en_25122825.mp3	United States English	Luther Rice College and Seminary is a private ...	2.0	0.0			{'array': [-9.094947017729282e-13, -9.094947017...]}
7676	common_voice_en_25122826.mp3	United States English	Songs from these albums have also been feature...	2.0	1.0			{'array': [-2.7284841053187847e-12, -3.6379788...]}
7677	common_voice_en_25122827.mp3	United States English	Another head had intake and exhaust ports reve...	2.0	0.0			{'array': [-2.2737367544323206e-11, 1.182343112...]}

7678 rows x 8 columns



In []:

```
set(data['accents'])
```

Out[]:

```
{'Australian English',  
 'Canadian English',  
 'England English',  
 'Filipino',  
 'India and South Asia (India, Pakistan, Sri Lanka)',  
 'Irish English',  
 'Scottish English',  
 'United States English'}
```

In []:

```
data['audio'][0]
```

Out[]:

```
{'array': array([-1.74622983e-10, -4.22005542e-10, -2.03726813e-10, ...,  
                4.54552310e-06, 1.05218151e-05, 1.69793075e-05]),  
 'path': './accents/datasets/common_voice_16_1/en/clips/common_voice_en_23608720.mp3',  
 'sampling_rate': 16000}
```

Data Pre-processing

Dealing with the class imbalance issues as the proportion of certain accents is far fewer than others. We first drop the columns that are unrelated to this research.

In []:

```
data = data.drop(['path'], axis=1)
data = data.drop(['up_votes'], axis=1)
data = data.drop(['down_votes'], axis=1)
data = data.drop(['age'], axis=1)
data = data.drop(['gender'], axis=1)
data['accents'].value_counts()
```

Out[]:

```
accents
United States English          3247
India and South Asia (India, Pakistan, Sri Lanka)  1188
England English                1094
Australian English             663
Filipino                      607
Canadian English              495
Irish English                  265
Scottish English              119
Name: count, dtype: int64
```

Then split the data to train and test set in the proportion of 8:2.

In []:

```
#train test split
train, test = train_test_split(data, test_size=0.2, random_state=42)

train['accents'].value_counts()
```

Out[]:

```
accents
United States English          2577
India and South Asia (India, Pakistan, Sri Lanka)  963
England English                891
Australian English             528
Filipino                      471
Canadian English              395
Irish English                  226
Scottish English              91
Name: count, dtype: int64
```

Then we upsample and downsample the minority and majority class.

In []:

```
import pandas as pd
from sklearn.utils import resample

# Down sampling three classes
US = train[train['accents'] == 'United States English']
US = US.sample(n=550)

India_south_asia = train[train['accents'] == 'India and South Asia (India, Pakistan, Sri Lanka)']
India_south_asia = India_south_asia.sample(n=550)

england = train[train['accents'] == 'England English']
england = england.sample(n=550)

# Upsample minority class
scottish = train[train['accents'] == 'Scottish English']
Irish = train[train['accents'] == 'Irish English']
scottish = resample(scottish, replace=True, n_samples= 270, random_state=42)
Irish = resample(Irish, replace=True, n_samples= 450, random_state=42)
```

```
# Combine all classes
the_rest = []
new = []
the_rest = pd.concat([train[train['accents'] == 'Australian English'], train[train['accents'] == "Filipino"], train[train['accents'] == 'Canadian English']])
new = pd.concat([US, India_south_asia, england, scottish, Irish, the_rest])

# Display new class counts
new.reset_index(drop=True, inplace=True)
new['accents'].value_counts()
```

Out[]:

```
accents
United States English      550
India and South Asia (India, Pakistan, Sri Lanka)  550
England English           550
Australian English         528
Filipino                   471
Irish English              450
Canadian English           395
Scottish English           270
Name: count, dtype: int64
```

In []:

```
train_ds = Dataset.from_pandas(new)
test_ds = Dataset.from_pandas(test)

train_ds = train_ds.class_encode_column("accents")
test_ds = test_ds.class_encode_column("accents")
```

Audio is not the interest of text classification, therefore remove the column.

In []:

```
train_ds = train_ds.remove_columns("audio")
test_ds = test_ds.remove_columns("audio")
```

In []:

```
labels = train_ds.features["accents"].names
label2id, id2label = dict(), dict()
for i, label in enumerate(labels):
    label2id[label] = str(i)
    id2label[str(i)] = label
```

In []:

```
train_ds
```

Out[]:

```
Dataset({
  features: ['accents', 'sentence'],
  num_rows: 3764
})
```

Bert

Tokenize the data

In []:

```
# tokenize the data
from transformers import AutoTokenizer, DataCollatorWithPadding

checkpoint = "bert-base-uncased"
```



```
def tokenize_function(example):
    return tokenizer(example["sentence"], padding = "max_length")

tokenized_datasets_train = train_ds.map(tokenize_function, batched=True)
tokenized_datasets_test = test_ds.map(tokenize_function, batched=True)

tokenized_datasets_train = tokenized_datasets_train.rename_column("accents", "label")
tokenized_datasets_test = tokenized_datasets_test.rename_column("accents", "label")

data_collator = DataCollatorWithPadding(tokenizer=tokenizer)
```

In []:

Out[]:

In []:

In []:

Define and train the BERT model with 8 epochs.

In []:

```
#Define a training arguments class
from transformers import TrainingArguments
training_args = TrainingArguments("test-trainer", evaluation_strategy="epoch", num_train_epochs=8)

#Define the model
num_labels = len(id2label)
num_labels
from transformers import AutoModelForSequenceClassification
model = AutoModelForSequenceClassification.from_pretrained(checkpoint,num_labels=num_labels)

#Define the compute metrics
import evaluate

def compute_metrics(eval_preds):
    metric = evaluate.load("accuracy")
    logits, labels = eval_preds
```

```
predictions = np.argmax(logits, axis=-1)
return metric.compute(predictions=predictions, references=labels)

#Define a trainer
from transformers import Trainer

trainer = Trainer(
    model,
    training_args,
    train_dataset=tokenized_datasets_train,
    eval_dataset=tokenized_datasets_test,
    data_collator=data_collator,
    tokenizer=tokenizer,
    compute_metrics = compute_metrics,
)
```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In []:

```
trainer.train()
```



Epoch	Training Loss	Validation Loss	Accuracy
1 No log		1.984597	0.160807
2 2.062500		1.947822	0.218099
3 1.762400		2.161429	0.187500
4 1.226100		2.723784	0.177083
5 0.585400		4.005829	0.184896
6 0.166800		4.573924	0.234375
7 0.049300		5.164403	0.202474
8 0.006100		5.232408	0.208984

Out[]:

TrainOutput(global_step=3768, training_loss=0.7775677567043375, metrics={'train_runtime': 875.9845, 'train_samples_per_second': 34.375, 'train_steps_per_second': 4.301, 'total_flos': 7923226912751616.0, 'train_loss': 0.7775677567043375, 'epoch': 8.0})

The plot on the changes of trainingh loss, validation loss and accuracy

In []:

```
import matplotlib.pyplot as plt

log_history = trainer.state.log_history

epochs = []
training_loss = [None] * len(epochs) # Initialize with None
validation_loss = [None] * len(epochs) # Initialize with None
accuracy = []

# Extracting data from log_history
for log in log_history:
    if 'epoch' in log:
        epoch_index = int(log['epoch']) - 1 # Convert to 0-indexed
        while len(epochs) <= epoch_index: # Append to epochs list if this epoch is not yet included
```

```

epochs.append(len(epochs) + 1)
training_loss.append(None) # Append None for this new epoch
validation_loss.append(None) # Append None for this new epoch
# Now we safely assume that the lengths of the lists are equal to epoch_index + 1
if 'loss' in log:
    training_loss[epoch_index] = log['loss']
if 'eval_loss' in log:
    validation_loss[epoch_index] = log['eval_loss']
if 'eval_accuracy' in log:
    accuracy.append(log['eval_accuracy'])

# Filter out the epochs where we have validation loss to match with training loss
valid_epochs = [epoch for epoch, val_loss in zip(epochs, validation_loss) if val_loss is
not None]
valid_validation_loss = [val_loss for val_loss in validation_loss if val_loss is not None]


# Comparison plot of training and validation loss
plt.figure(figsize=(8, 4))
plt.plot(epochs, training_loss, label='Training Loss', marker='o')
plt.plot(valid_epochs, valid_validation_loss, label='Validation Loss', marker='x')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training vs Validation Loss')
plt.legend()
plt.grid(True)
plt.show()

# Ensure the list lengths match for accuracy plot
while len(accuracy) < len(epochs):
    accuracy.append(None)


# Separate plot for accuracy
plt.figure(figsize=(8, 4))
plt.plot(epochs, accuracy, label='Accuracy', color='green', marker='o')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Accuracy per Epoch')
plt.legend()
plt.grid(True)
plt.show()


```

Training vs Validation Loss



Accuracy per Epoch





Confusion Matrix on the BERT Transformer

In []:

```
#confusion matrix
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

predictions = trainer.predict(tokenized_datasets_test)
y_pred = np.argmax(predictions.predictions, axis=1)
y_true = tokenized_datasets_test['label']
cm = confusion_matrix(y_true, y_pred)

plt.figure(figsize=(5, 4))
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix with True and False Labels')
ax = sns.heatmap(cm, annot=True, fmt='g', cmap='Blues', xticklabels=id2label.values(), yticklabels=id2label.values())
```

Confusion Matrix with True and False Labels

	Australian English	Canadian English	England English	Filipino	India and South Asia (India, Pakistan, Sri Lanka)	Irish English	Scottish English	United States English
Australian English	29	3	28	19	25	0	1	30
Canadian English	18	19	19	14	9	0	0	21
England English	45	11	40	35	18	1	3	50
Filipino	21	5	39	23	16	1	0	31
India and South Asia (India, Pakistan, Sri Lanka)	27	9	41	41	46	1	0	60
Irish English	9	1	9	8	4	0	0	8
Scottish English	8	1	11	1	1	0	0	6
United States English	153	44	110	110	85	2	2	164

5 epoch

Define and train the model with 5 epochs.

In []:

```
#Define a training arguments class
from transformers import TrainingArguments
training_args = TrainingArguments("test-trainer", evaluation_strategy="epoch", num_train_epochs=5)

#Define the model
num_labels = len(id2label)
num_labels
from transformers import AutoModelForSequenceClassification
model = AutoModelForSequenceClassification.from_pretrained(checkpoint, num_labels=num_labels)

#Define the compute metrics
import evaluate

def compute_metrics(eval_preds):
    metric = evaluate.load("accuracy")
    logits, labels = eval_preds
    predictions = np.argmax(logits, axis=-1)
    return metric.compute(predictions=predictions, references=labels)

#Define a trainer
from transformers import Trainer

trainer = Trainer(
    model,
    training_args,
    train_dataset=tokenized_datasets_train,
    eval_dataset=tokenized_datasets_test,
    data_collator=data_collator,
    tokenizer=tokenizer,
    compute_metrics = compute_metrics,
)
```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In []:

```
trainer.train()
```

[2355/2355 09:07, Epoch 5/5]

Epoch	Training Loss	Validation Loss	Accuracy
1 No log		2.004365	0.134766
2 2.071700		1.953544	0.162109

3	1.789400		2.034272	0.203776
Epoch	Training Loss	Validation Loss	Accuracy	
4	1.268700	2.584999	0.171224	
5	0.686900	2.747876	0.225260	

Out[]:

```
TrainOutput(global_step=2355, training_loss=1.2882614556897725, metrics={'train_runtime': 548.8262, 'train_samples_per_second': 34.291, 'train_steps_per_second': 4.291, 'total_flos': 4952016820469760.0, 'train_loss': 1.2882614556897725, 'epoch': 5.0})
```

Do the plotting as above.

In []:

```
import matplotlib.pyplot as plt

# Assuming trainer.state.log_history contains the logged metrics
log_history = trainer.state.log_history

epochs = []
training_loss = [None] * len(epochs) # Initialize with None
validation_loss = [None] * len(epochs) # Initialize with None
accuracy = []

# Extracting data from log_history
for log in log_history:
    if 'epoch' in log:
        epoch_index = int(log['epoch']) - 1 # Convert to 0-indexed
        while len(epochs) <= epoch_index: # Append to epochs list if this epoch is not yet included
            epochs.append(len(epochs) + 1)
            training_loss.append(None) # Append None for this new epoch
            validation_loss.append(None) # Append None for this new epoch
        # Now we safely assume that the lengths of the lists are equal to epoch_index + 1
        if 'loss' in log:
            training_loss[epoch_index] = log['loss']
        if 'eval_loss' in log:
            validation_loss[epoch_index] = log['eval_loss']
        if 'eval_accuracy' in log:
            accuracy.append(log['eval_accuracy'])


# Filter out the epochs where we have validation loss to match with training loss
valid_epochs = [epoch for epoch, val_loss in zip(epochs, validation_loss) if val_loss is not None]
valid_validation_loss = [val_loss for val_loss in validation_loss if val_loss is not None]

# Comparison plot of training and validation loss
plt.figure(figsize=(8, 4))
plt.plot(epochs, training_loss, label='Training Loss', marker='o')
plt.plot(valid_epochs, valid_validation_loss, label='Validation Loss', marker='x')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training vs Validation Loss')
plt.legend()
plt.grid(True)
plt.show()


# Ensure the list lengths match for accuracy plot
while len(accuracy) < len(epochs):
    accuracy.append(None)

# Separate plot for accuracy
plt.figure(figsize=(8, 4))
plt.plot(epochs, accuracy, label='Accuracy', color='green', marker='o')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Accuracy per Epoch')
plt.legend()
plt.grid(True)
plt.show()
```

Training vs Validation Loss



Accuracy per Epoch



And this is the confusion matrix.

In []:

```
#confusion matrix
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

predictions = trainer.predict(tokenized_datasets_test)
y_pred = np.argmax(predictions.predictions, axis=1)
y_true = tokenized_datasets_test['label']
cm = confusion_matrix(y_true, y_pred)

plt.figure(figsize=(5, 4))
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix with True and False Labels')
ax = sns.heatmap(cm, annot=True, fmt='g', cmap='Blues', xticklabels=id2label.values(), yticklabels=id2label.values())
```

Confusion Matrix with True and False Labels

Confusion Matrix with True and False Labels

